

ASHIS SHARMA

📍 Hyderabad ✉ ashishsharma@outlook.com ☎ +918759150942 🌐 in/ashishsharma 🖱 ashishsharma.com

SUMMARY

DevOps & AI Infrastructure Engineer with 5+ years' experience scaling GPU-based LLM training clusters (100+ GPUs) and automating multi-cloud MLOps pipelines across GCP and Azure. Proven success reducing deployment times by 60% and optimizing GPU utilization through Terraform, Helm, and Kubeflow. Passionate about building scalable AI infrastructure that bridges model research and production reliability.

EXPERIENCE

Jukshio

DevOps Engineer

Hyderabad, India
June 2020- Present

- Configured and optimized Slurm-based GPU clusters for large-scale LLM fine-tuning (100+ GPUs).
- Deployed scalable Kubernetes inference services using SGLang and vLLM, improving latency and throughput for customer workloads.
- Automated infrastructure provisioning with Terraform, Helm, and Kustomize, reducing setup time by 40%.
- Designed multi-stage CI/CD pipelines (build, test, security scan, deployment) in GitLab CI, cutting release cycles by 60%.
- Implemented DevSecOps practices with SonarQube, Trivy, and OWASP WAF, enhancing compliance and security posture.
- Built and automated Kubeflow-based MLOps pipeline supporting 20+ model training jobs across GPU clusters.
- Managed Prometheus and Grafana observability stack; reduced time-to-detect incidents by 40%.
- Migrated workloads from Azure to GCP, improving reliability and reducing cloud costs by 25%.
- Developed SDKs for enterprise clients (Jio, HDFC Bank), streamlining integration workflows.
- Created internal dashboards in React + Node.js to automate reporting and enhance productivity.

Wipro

Test Engineer

Chennai, India
August 2019 - March 2020

- Automated regression testing with Selenium WebDriver (Java) and built data validation utilities to accelerate QA cycles.

PROJECT

GPU Cluster Orchestration for LLM Training

- Configured and automated large-scale Slurm-based GPU clusters (100+ GPUs) to support fine-tuning and inference of open-source LLMs.
- Implemented auto-scaling and resource optimization, improving GPU utilization and cutting idle cloud costs by 25%.

End-to-End MLOps Pipeline with Kubeflow

- Built and maintained a Kubeflow-based pipeline enabling reproducible model training, validation, and deployment across GPU clusters.
- Reduced model iteration cycle time by 35% through workflow automation and CI/CD integration.

Edge IoT Cluster Management System

- Configured and managed K3s clusters for SmartCart IoT devices deployed in remote locations.
- Enabled remote deployment, updates, and monitoring for 500+ edge devices through a centralized control plane.

CI/CD and DevSecOps Automation Framework

- Designed a multi-stage CI/CD pipeline (build, test, scan, deploy) with GitLab CI and Terraform, reducing deployment time by 60%.
- Integrated Trivy, SonarQube, and WAF to automate security scans and ensure ISO compliance.

EDUCATION

Bachelor of Science in Information Technology

SRM University Sikkim

Gangtok
2019

SKILLS

Cloud & Infrastructure: Google Cloud Platform (GCP), Microsoft Azure, Terraform
Containers & Orchestration: Kubernetes, Docker, Helm, Kustomize, K3s, Slurm
DevOps & Automation: GitLab CI/CD, GitHub Actions, Multi-stage Pipelines
Monitoring & Security: Prometheus, Grafana, Loki, EFK Stack, Trivy, SonarQube, WAF, OWASP Rules
Programming: GoLang, Python, JavaScript, Bash
AI & HPC: CUDA, TensorRT, vLLM, SGLang, Kubeflow